RAFT - Distributed consensus protocol

Distributed = Many
Consensus = Agreement

Consistent core

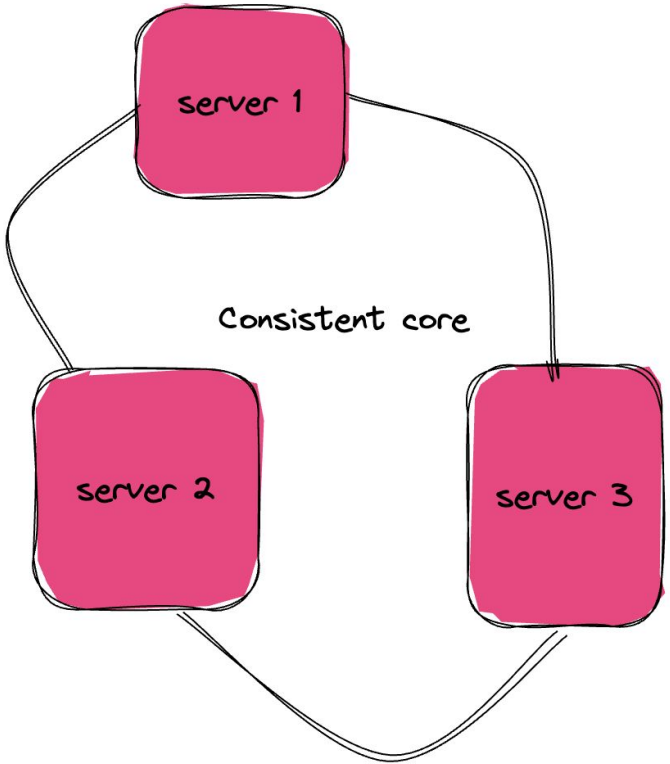

Linearizability

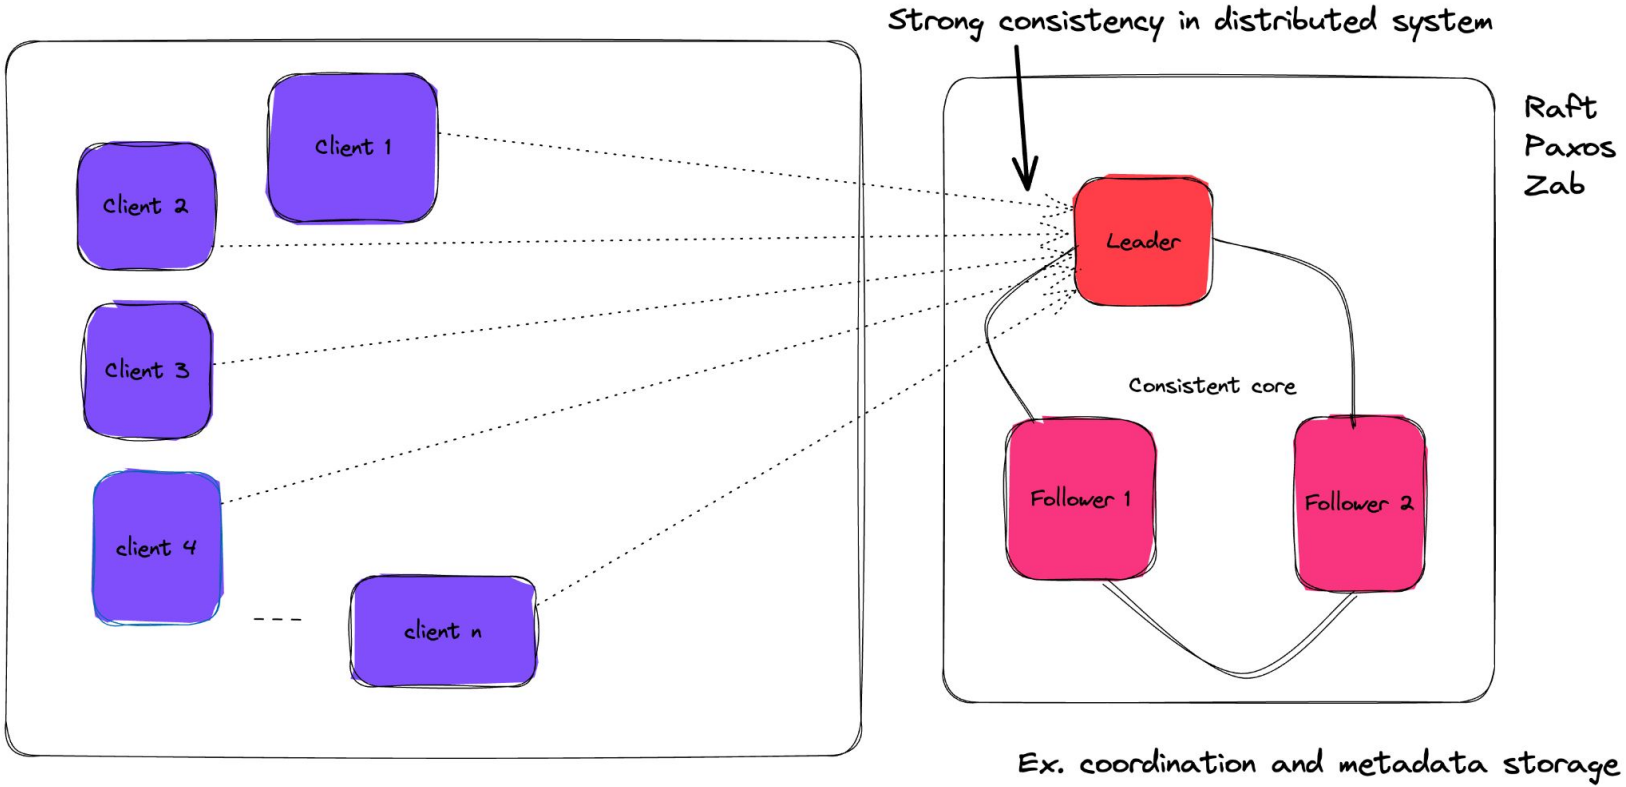

Examples

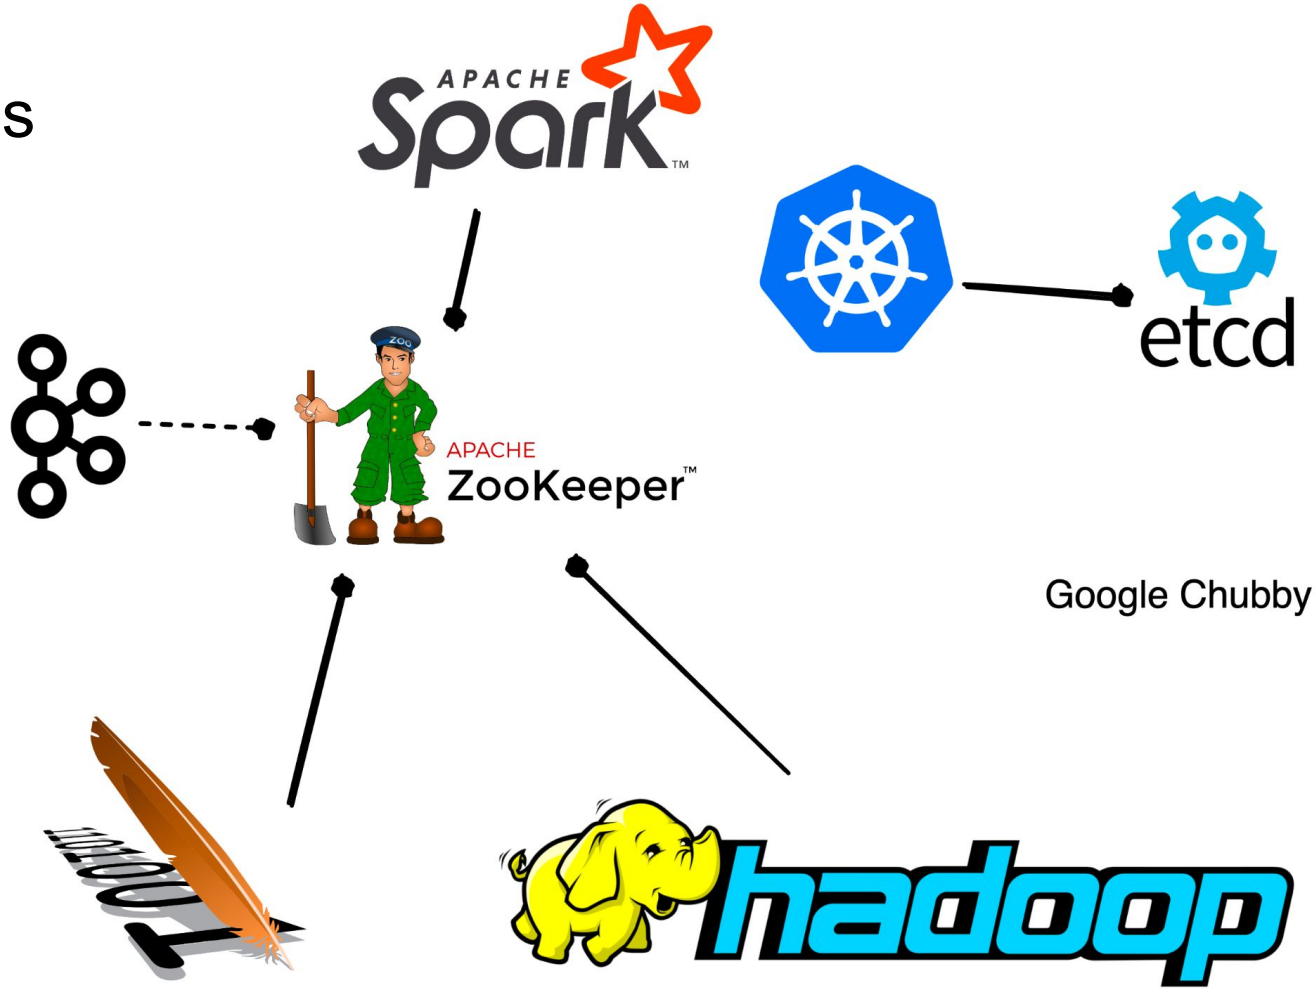

Paxos

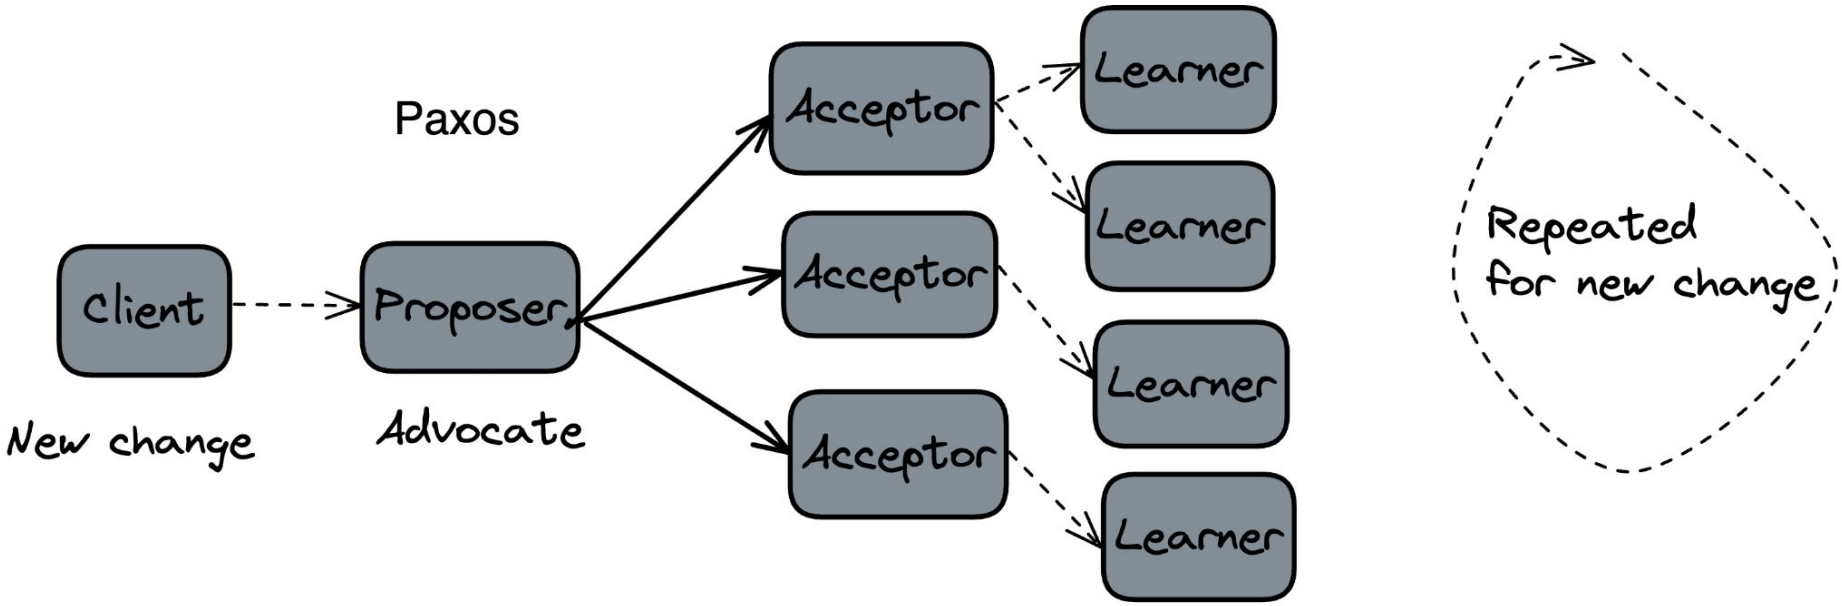
Paxos
Client
Proposer
Acceptor
Acceptor
Acceptor
Learner
Learner
Learner
Learner
New change
Advocate
Repeated
for new change

Roles in Raft

It has 3 roles

- Leader
- Follower
- Candidate

Example

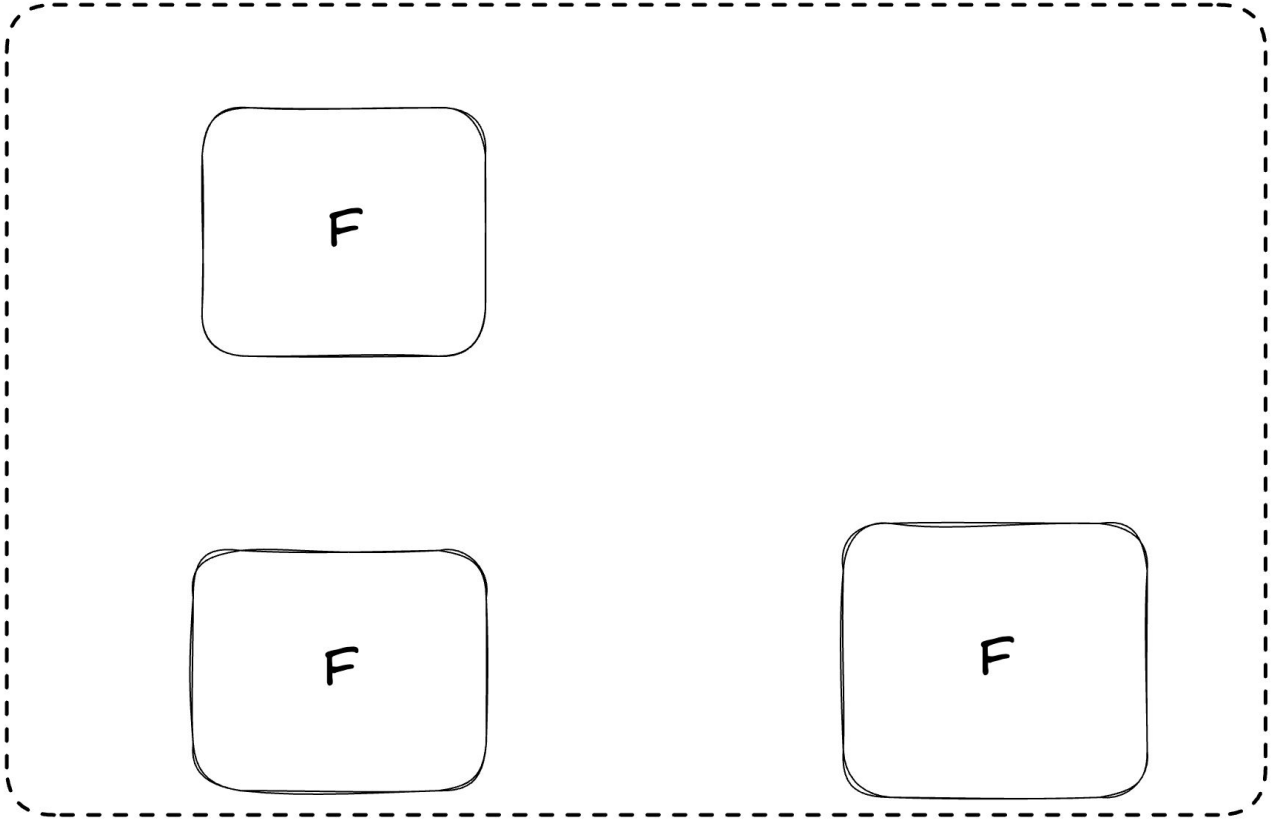

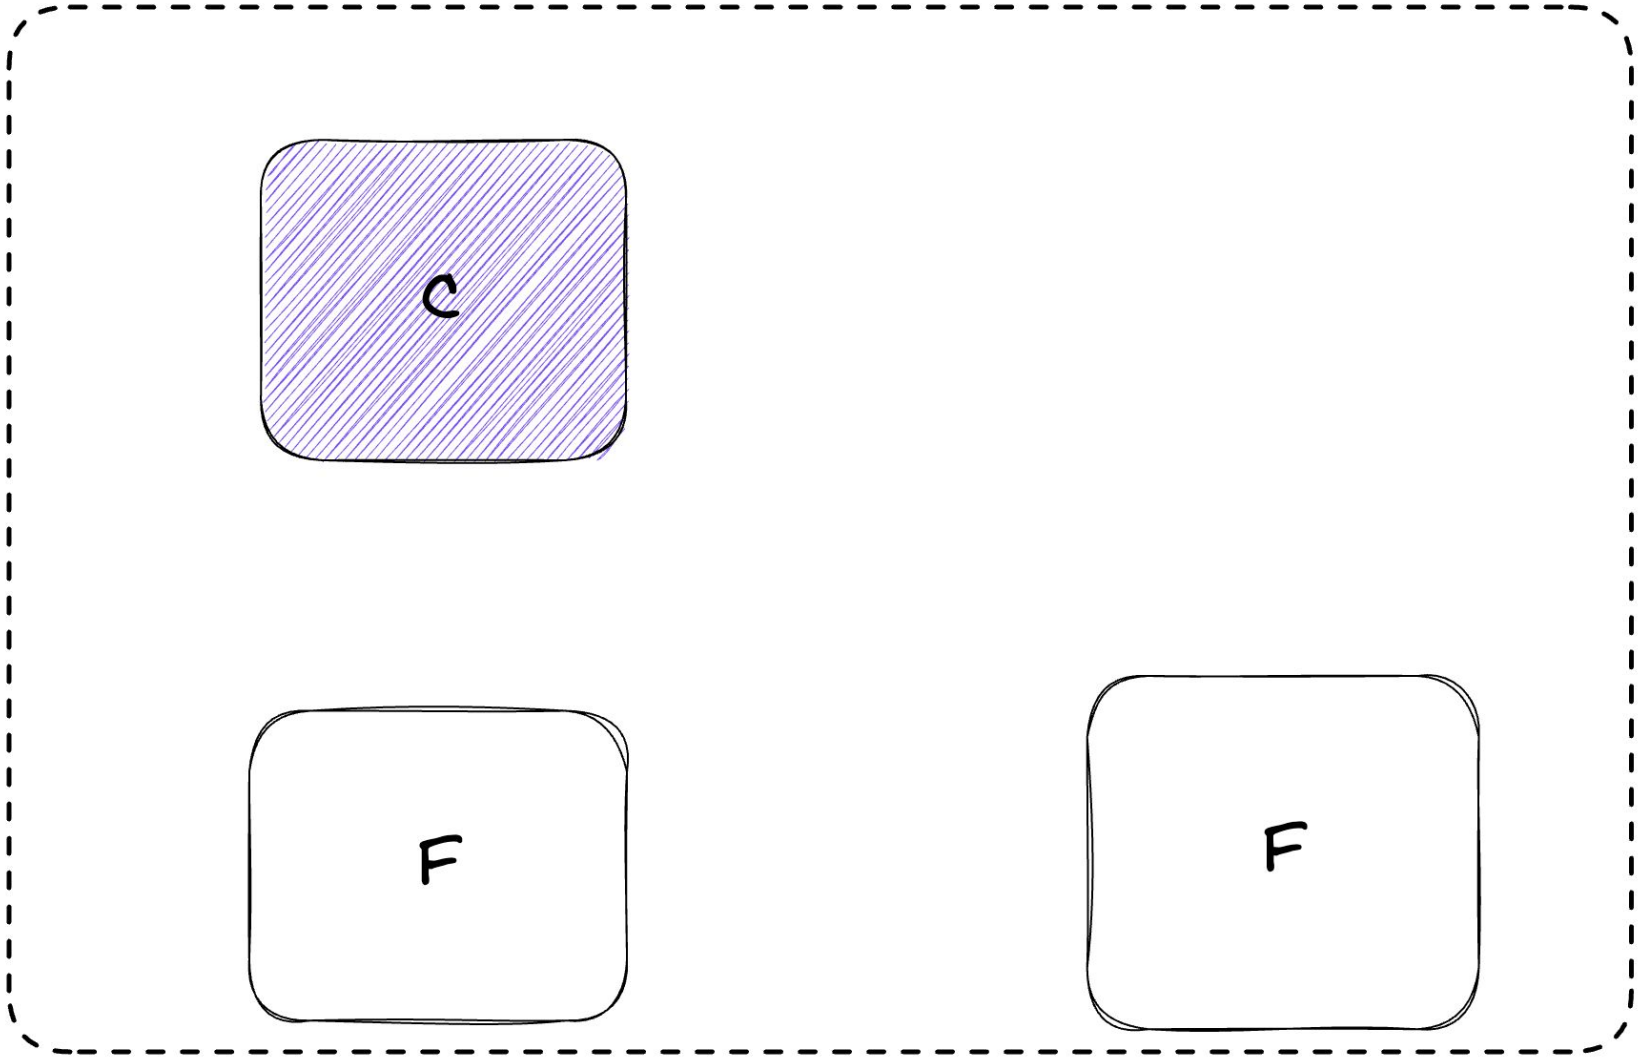
C
F
F

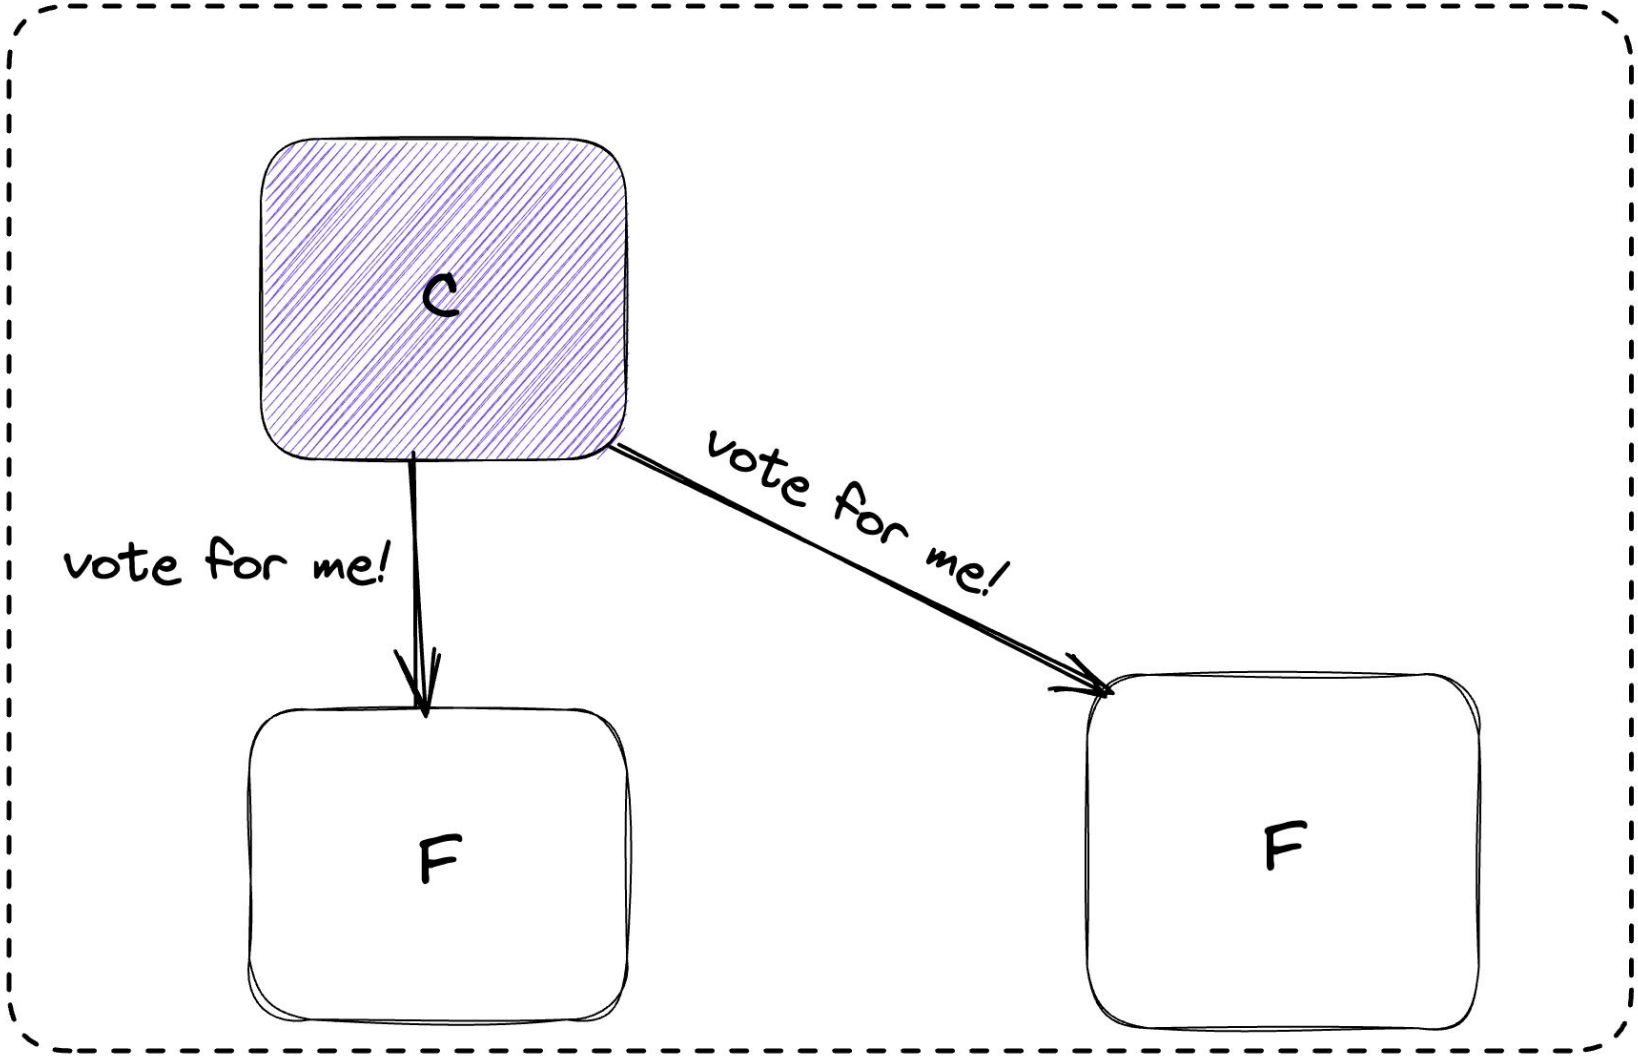
C
vote for me!
vote for me!
F
F

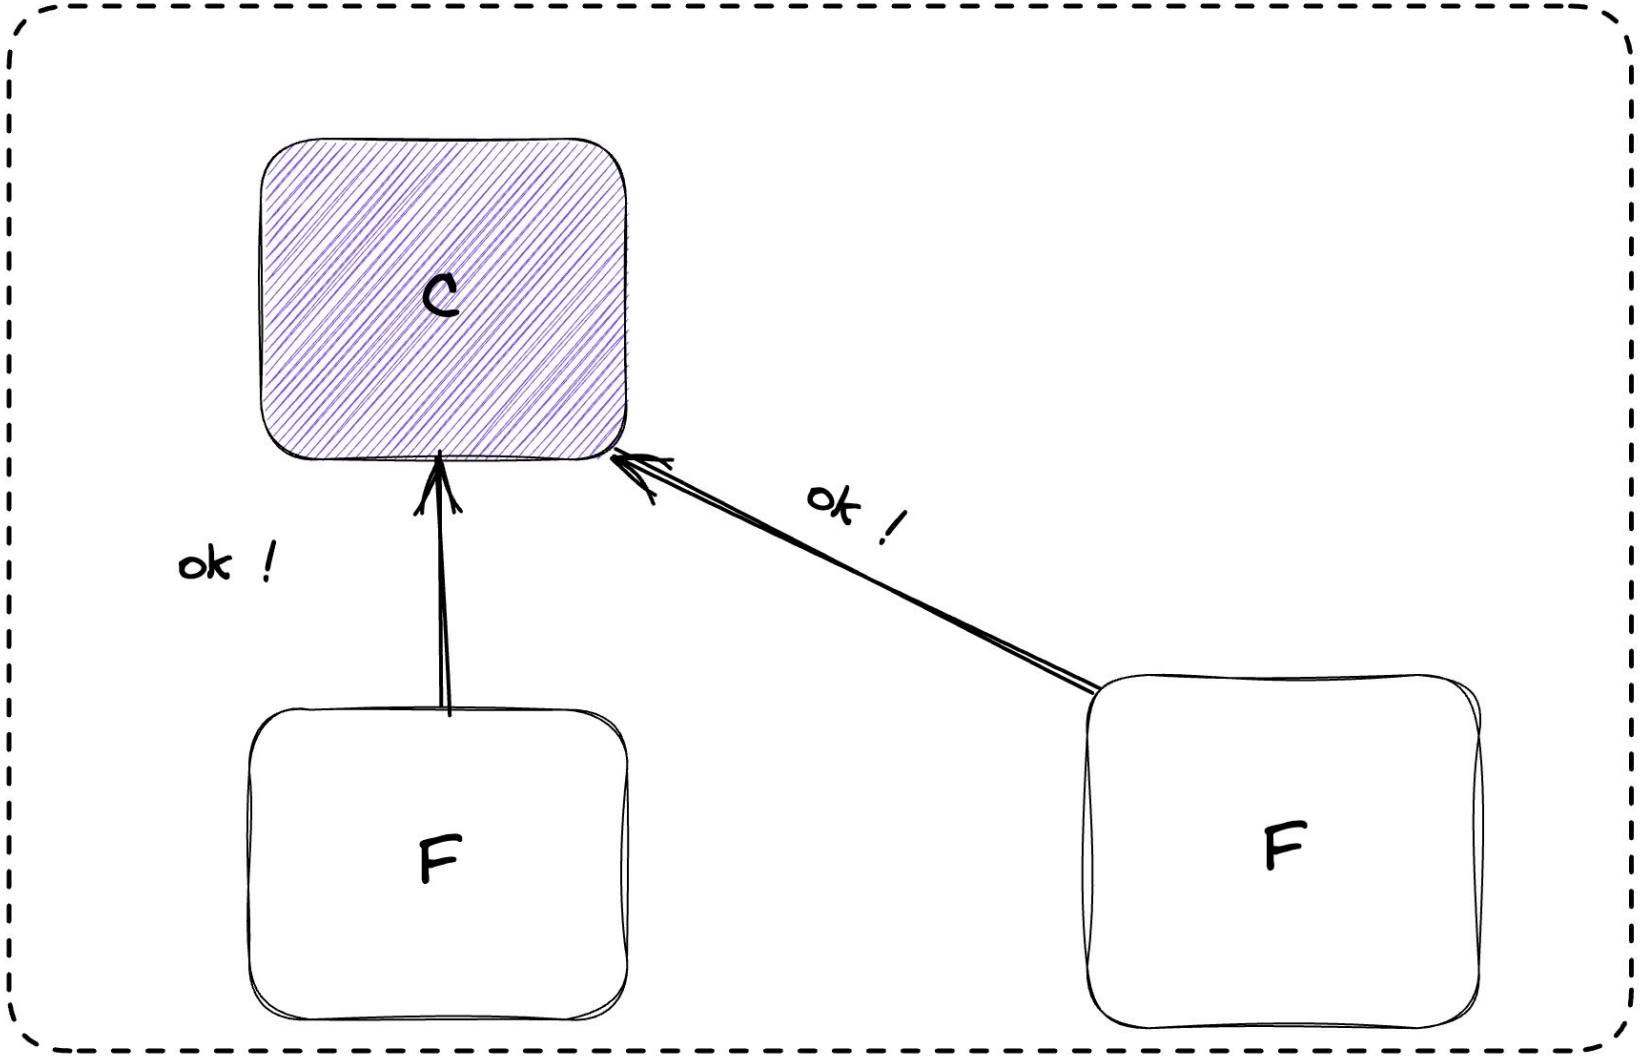
C
ok !
ok !
F
F

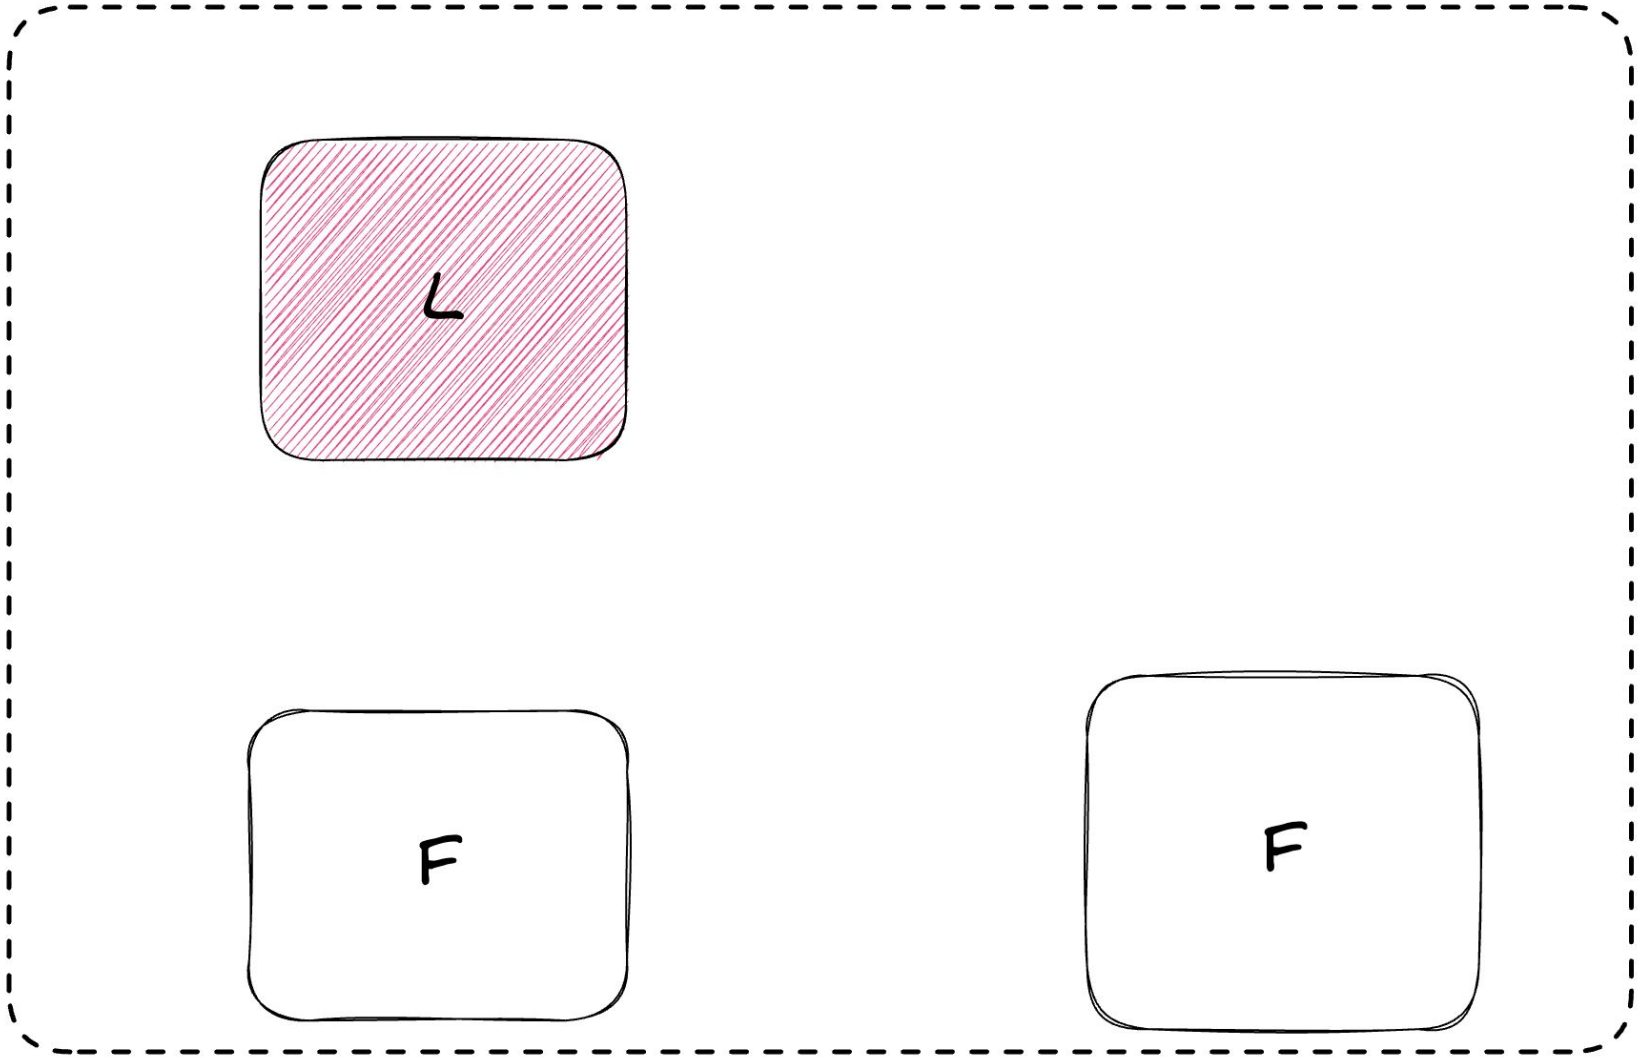
L
F
F

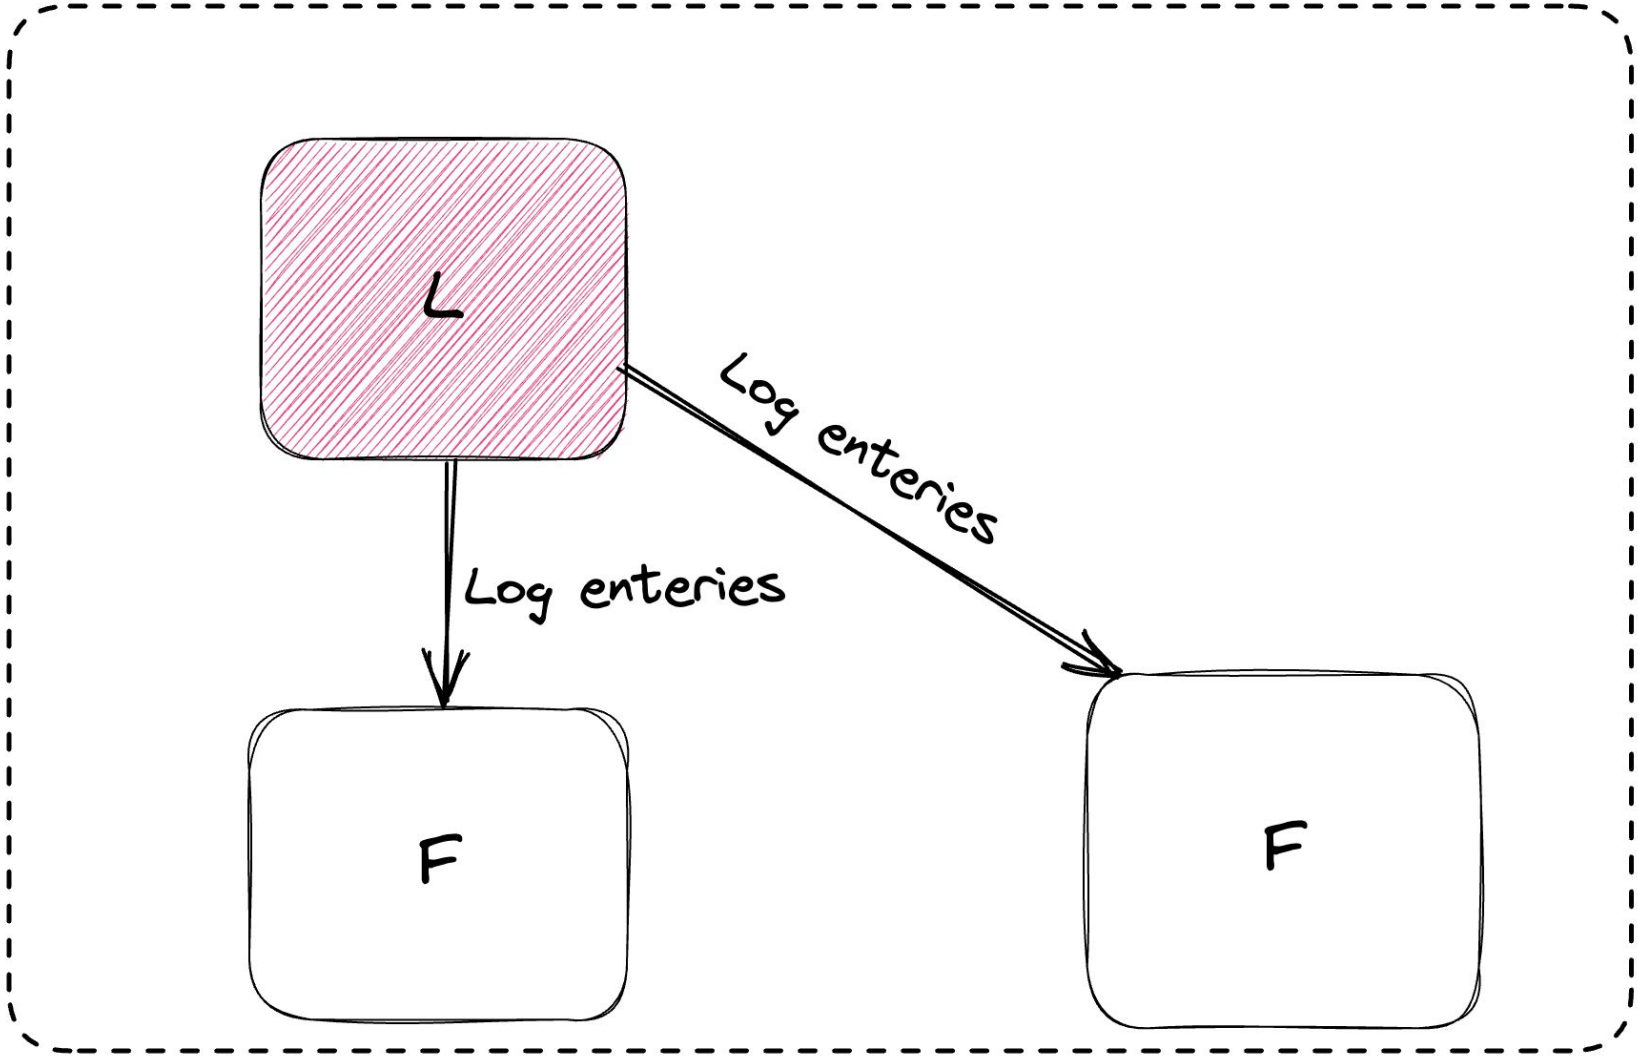
L
Log enteries
Log enteries
F
F

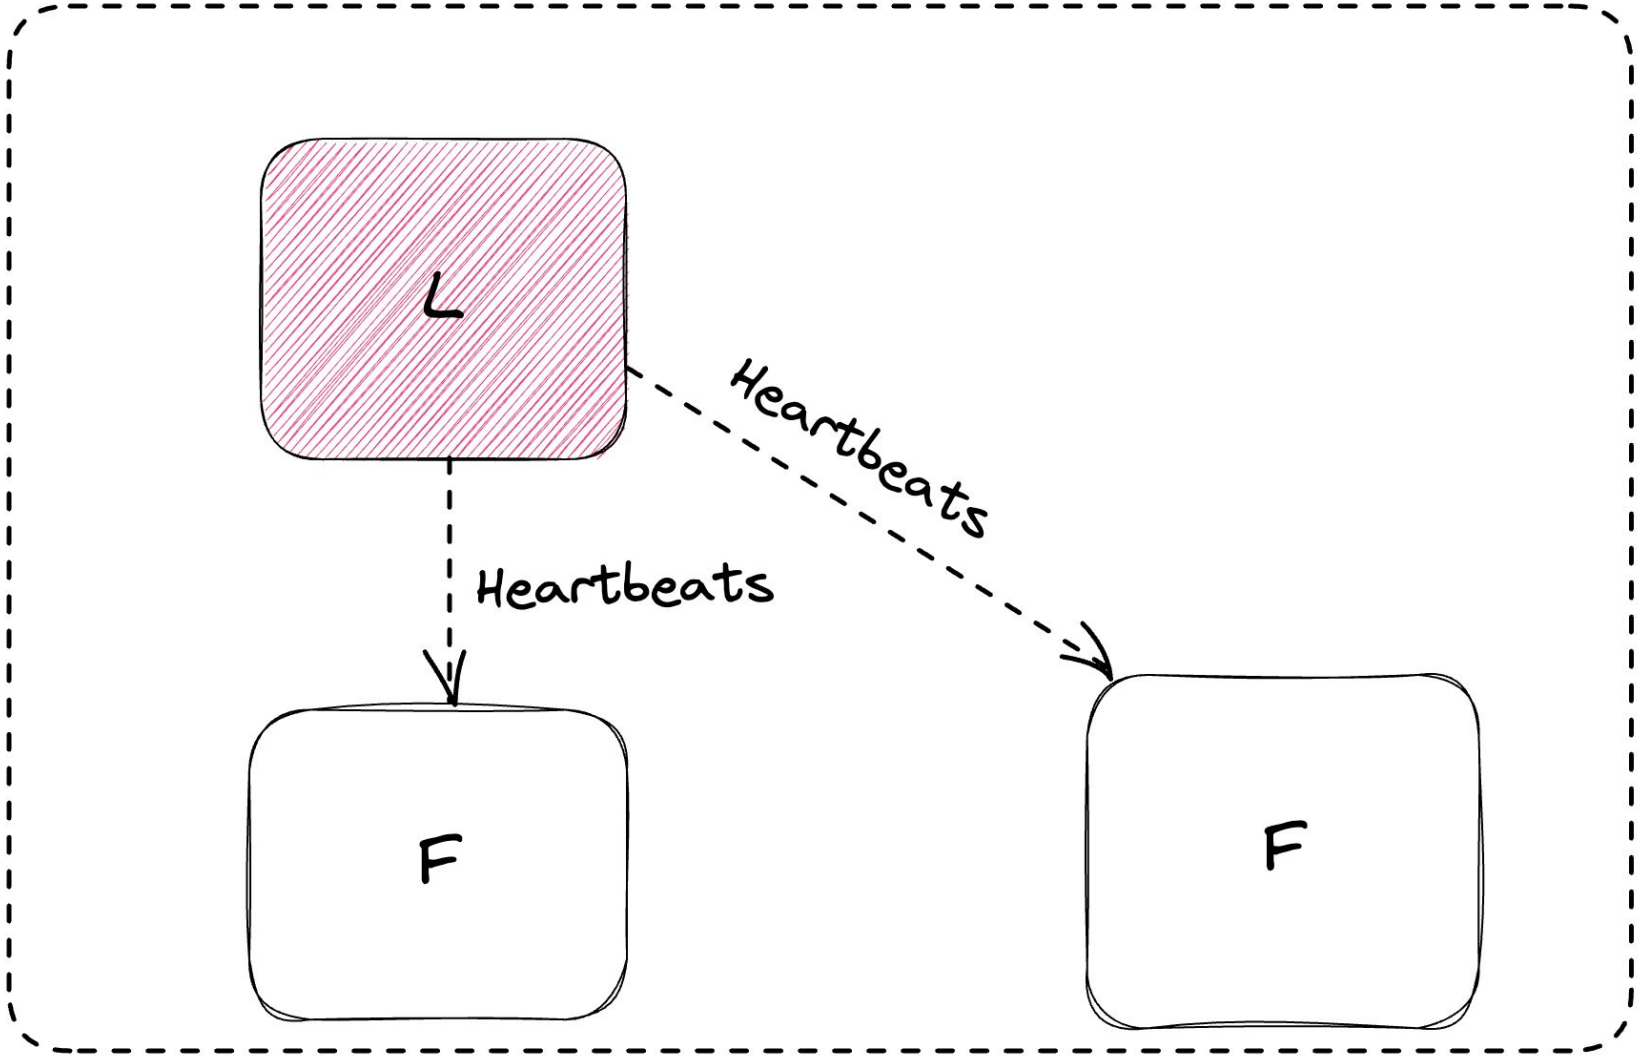
L
Heartbeats
Heartbeats
F
F

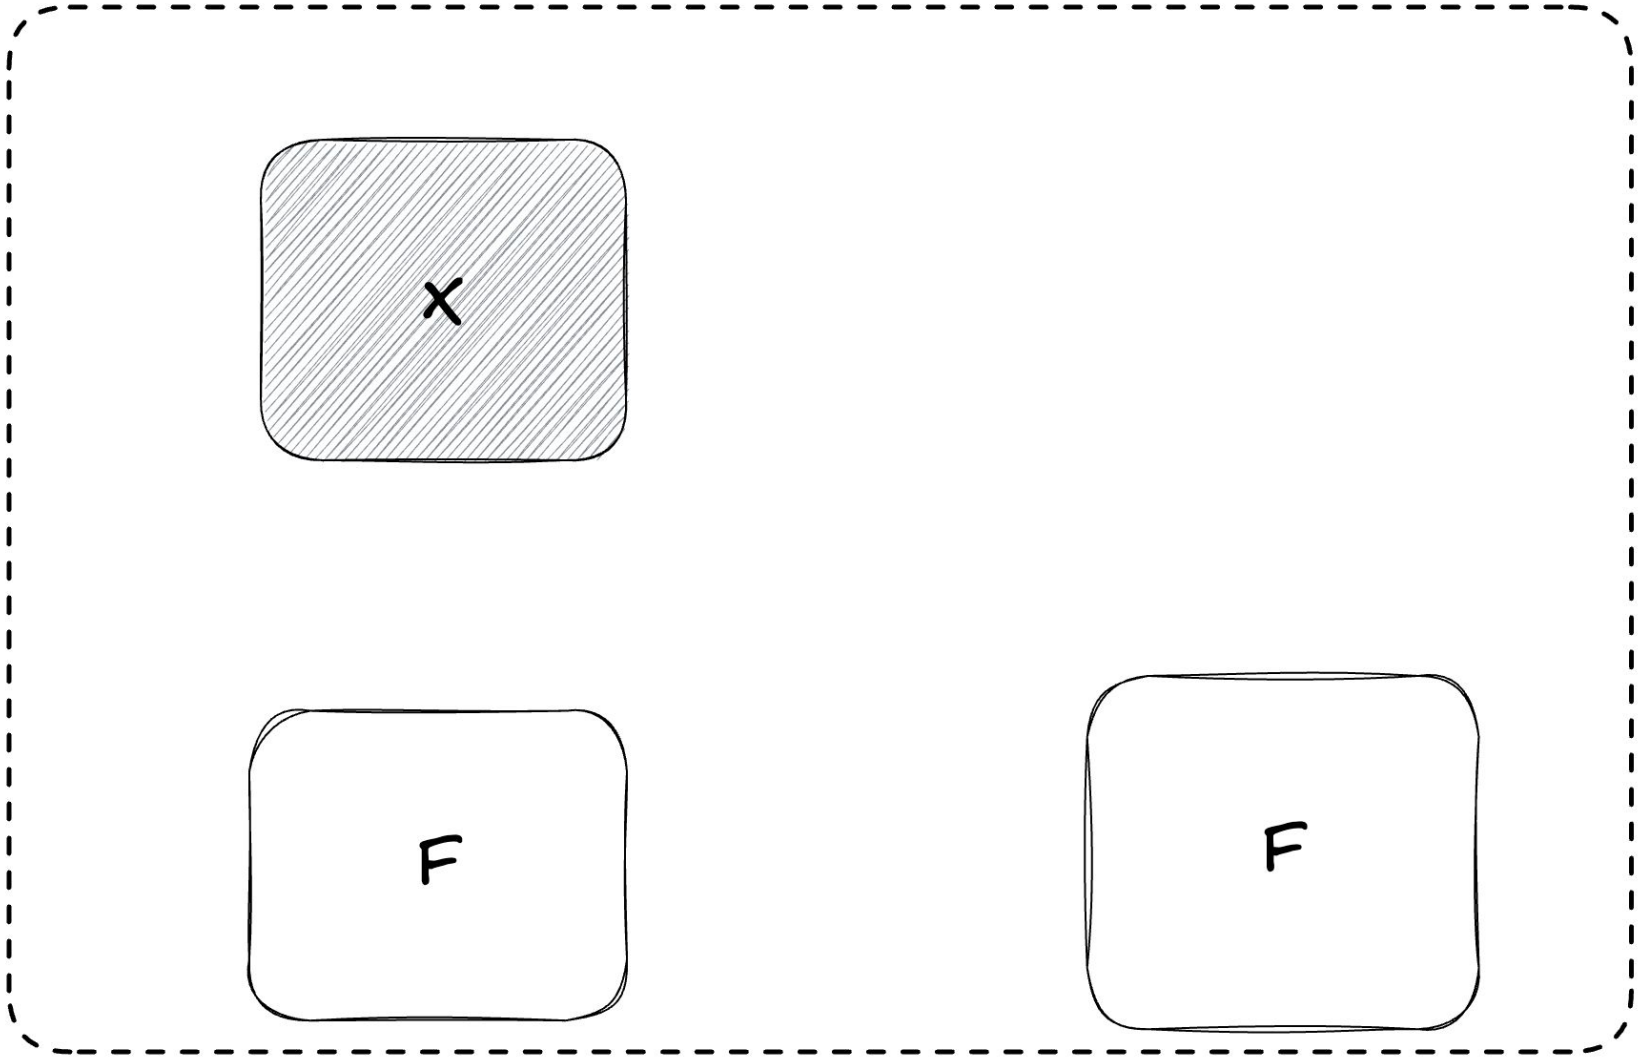
X
F
F

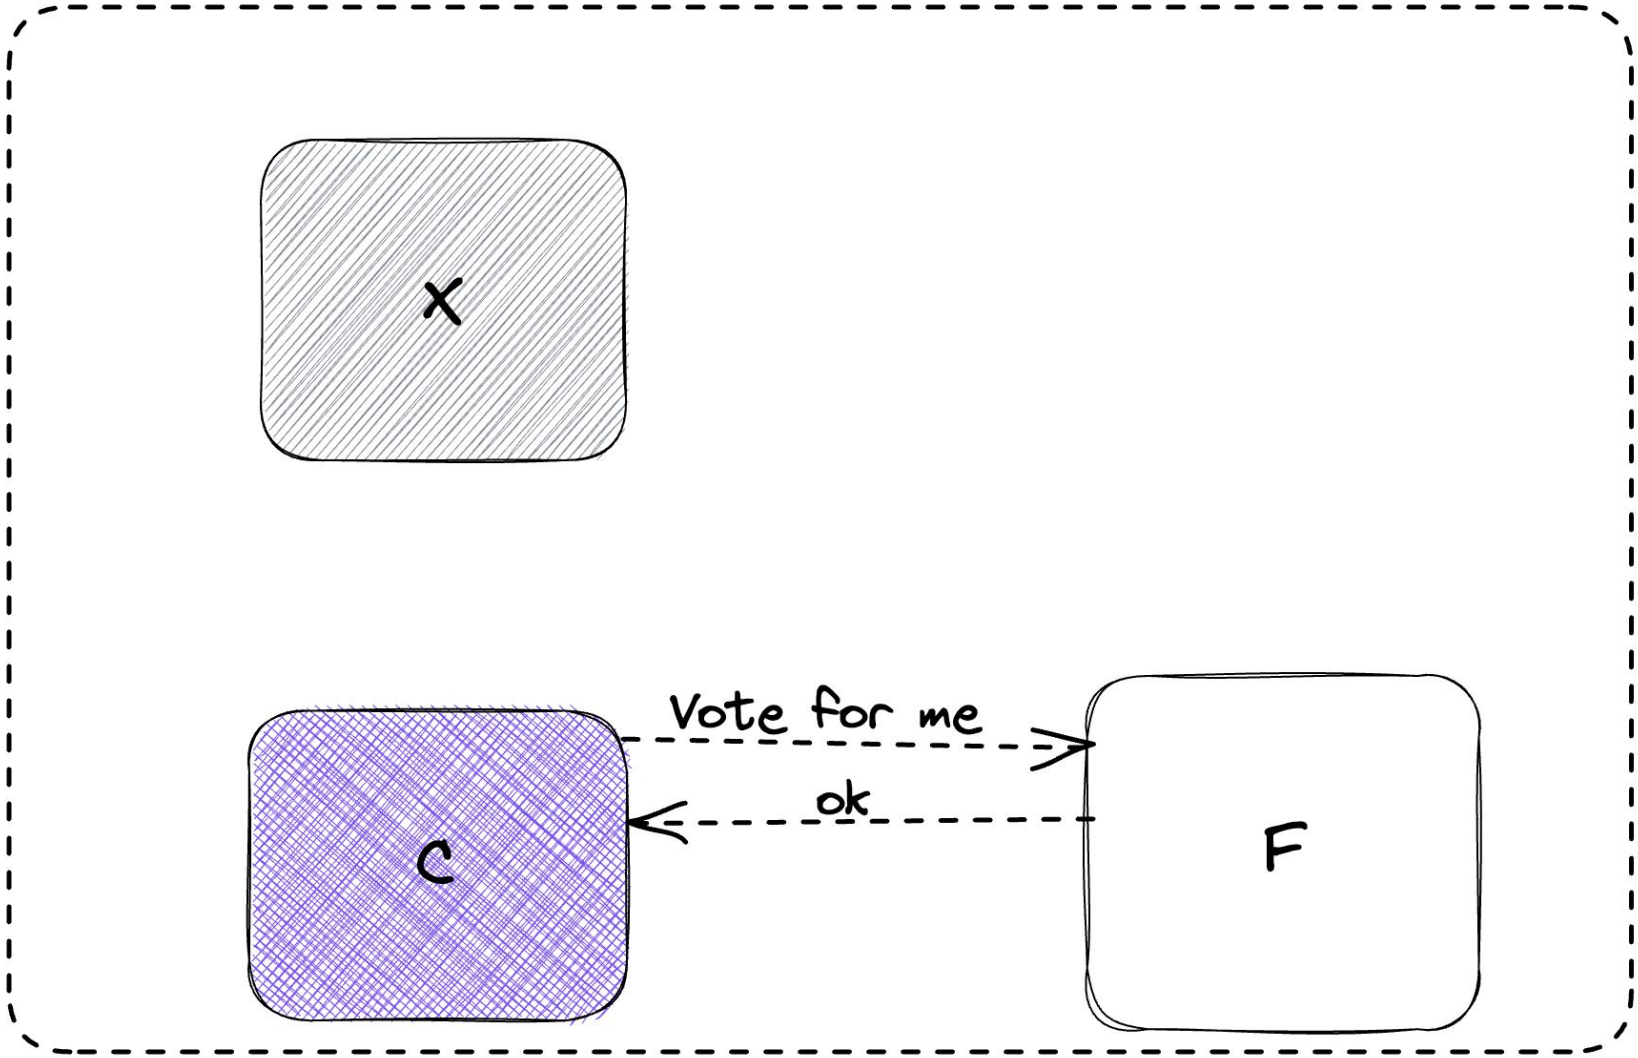
X
C
Vote for me
ok
F

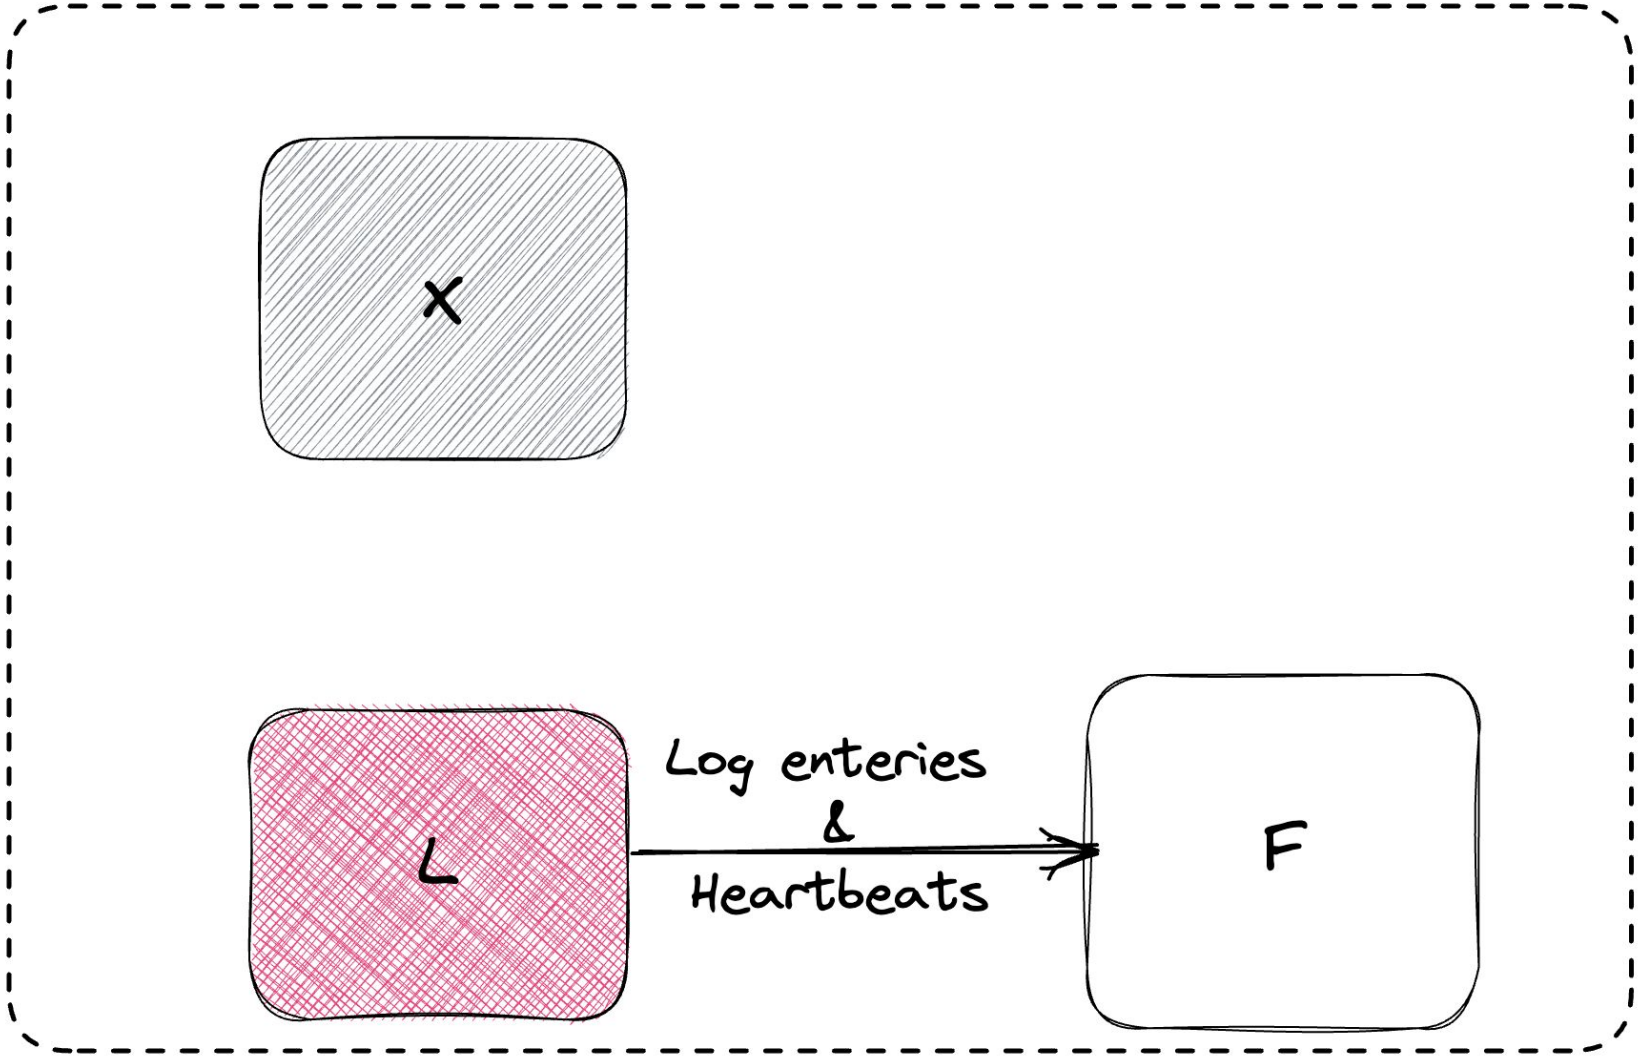
X
L
Log enteries
&
Heartbeats
F

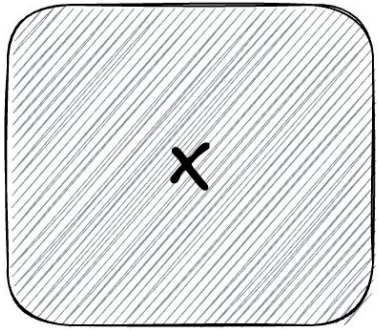

System is down as It need
more than 1 votes to
promote to leader in 3 node system

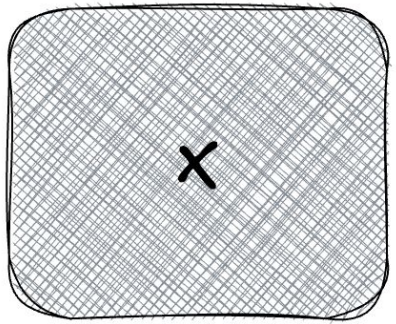

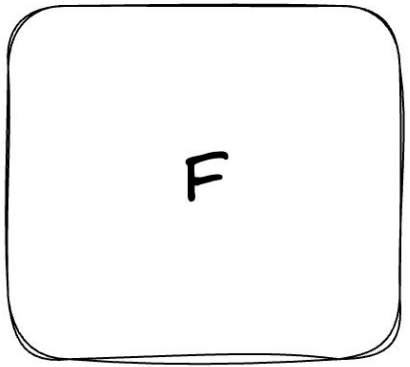

Leader Election (Best case scenario)

Refer diagram attached separately

Leader Election (Split Vote)

Refer diagram attached separately

Log Replication

Refer diagram attached separately

Log Replication (with network partitions)

Refer diagram attached separately

A Byzantine failure

a condition of a distributed computing systems, where **components may fail and there is imperfect information** on whether a component has failed.

a component such as a server can inconsistently **appear both failed and functioning** to failure-detection systems, presenting different symptoms to different observers.

etcd - A Byzantine failure is real

- When cluster leader fails, **etcd uses RAFT protocol** to maintain consistency and establish consensus to promote a new leader.
- In the RAFT protocol, cluster members are assumed to be either **available or unavailable**, and to provide **accurate information or none at all**.
- This works fine when a machine crashes, but is not always able to handle situations where different members of the cluster have **conflicting information**.

RAFT - A Byzantine failure is real

This caused cluster members to have conflicting views of reality, known in distributed systems theory as a Byzantine fault.

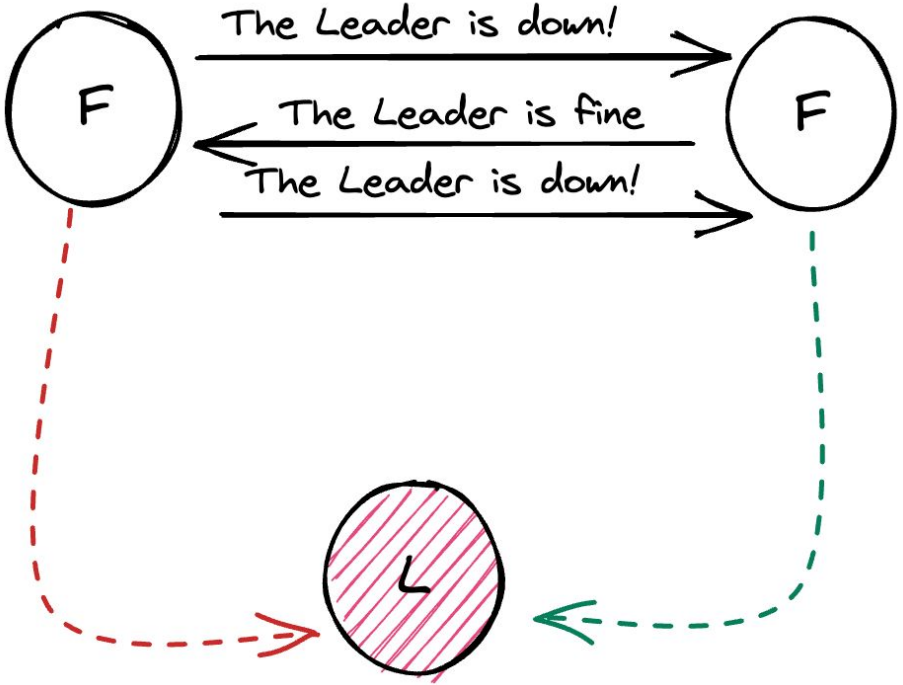

In 2020, Problem faced by

CLOUDFLARE

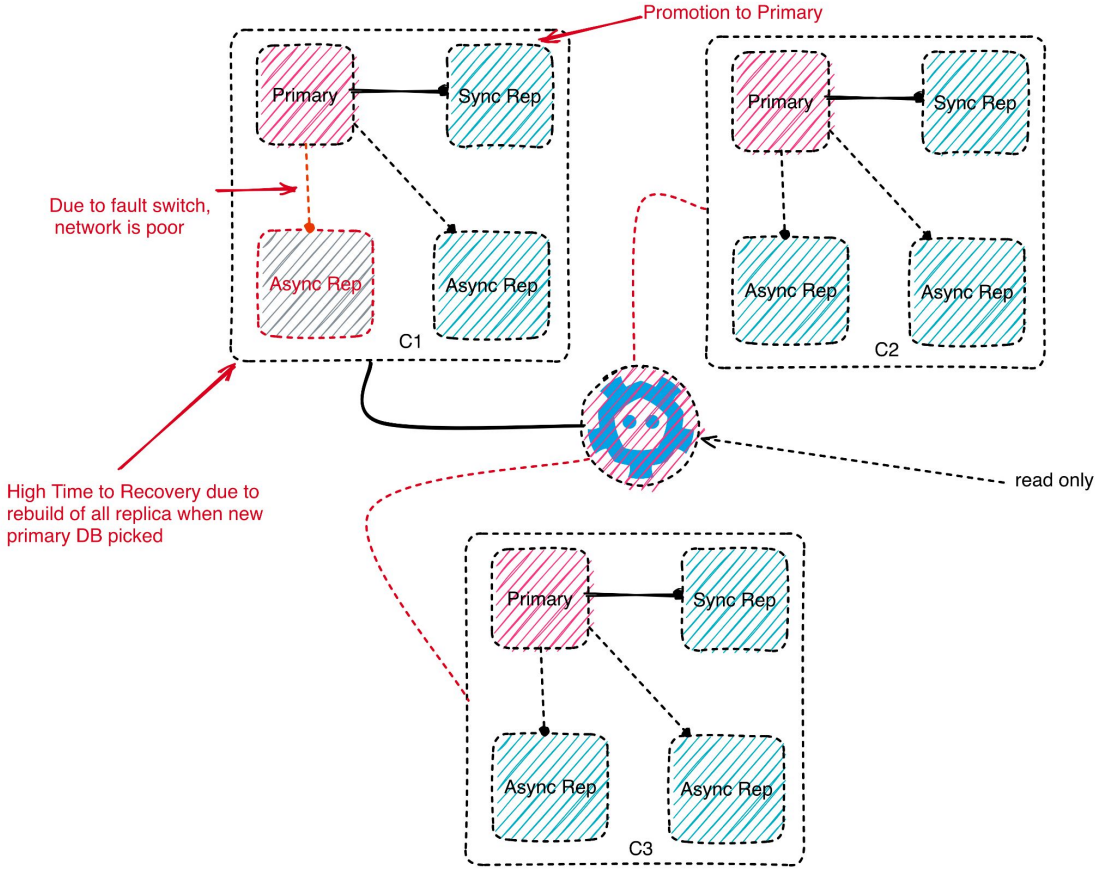
Promotion to Primary
Primary
Sync Rep
Async Rep
Async Rep
C1
Due to fault switch,
network is poor
Primary
Sync Rep
Async Rep
Async Rep
C2
read only
High Time to Recovery due to
rebuild of all replica when new
primary DB picked
Primary
Sync Rep
Async Rep
Async Rep
C3

Impact of Byzantine failure in cloudflare

- API success rate **dipped to 75%**.
- Redundant cluster nodes were unreachable due to **limbo in leader election**.
- **Read only** operation permitted to database.
- More **pressure to single cluster** for serving traffic.
- Downtime of around **6 hrs 33 min**.
- Dashboard api response **80 times slower** compared to normal.

Thank you

Happy Learning